

CS/DSC/AI 391L: Machine Learning**Homework 1 - Theory***Lecture: Prof. Adam Klivans**Keywords: Boolean functions, mistake bounds, PAC learning*

Instructions: Please either typeset your answers (L^AT_EX recommended) or write them very clearly and legibly and scan them, and upload the PDF on edX. Legibility and clarity are critical for fair grading.

- Often in binary classification we are interested in the differences in the output of our current classifier, g , and an unknown function f that we are trying to learn. It is common in these cases to examine the quantity produced by $f(x)g(x)$ for a given input x . For this problem, let D be an arbitrary distribution on the domain $\{-1, 1\}^n$, and let $f, g : \{-1, 1\}^n \rightarrow \{-1, 1\}$ be two Boolean functions.

- [6 points]** Prove that

$$\mathbb{P}_{x \sim D}[f(x) \neq g(x)] = \frac{1 - \mathbb{E}_{x \sim D}[f(x)g(x)]}{2}.$$

- [4 points]** Would this still be true if the domain were some other domain (such as \mathbb{R}^n , where \mathbb{R} denotes the real numbers, with say the Gaussian distribution) instead of $\{-1, 1\}^n$? If yes, justify your answer. If not, give a counterexample.

Note: Only the *domain* changes here. The output is still boolean.

- [10 points]** Let f be a decision tree with t leaves over the variables $x = (x_1, \dots, x_n) \in \{-1, 1\}^n$. Explain how to write f as a multivariate polynomial $p(x_1, \dots, x_n)$ such that for every input $x \in \{-1, 1\}^n$, $f(x) = p(x)$. (You may interpret -1 as FALSE and 1 as TRUE or the other way round, at your preference.) (*Hint: try to come up with an “indicator polynomial” for every leaf, i.e. one that evaluates to the leaf’s value if x is such that that path is taken, and 0 otherwise.*)
- [10 points]** Compute a depth-two decision tree for the training data in table 1 using the Gini function, $C(a) = 2a(1 - a)$ as described in class. What is the overall accuracy on the training data of the tree? For clarity, this will be a full binary tree and a full binary tree of depth-two has four leaves.

X	Y	Z	Number of positive examples	Number of negative examples
0	0	0	10	20
0	0	1	25	5
0	1	0	35	15
0	1	1	35	5
1	0	0	5	15
1	0	1	30	10
1	1	0	10	10
1	1	1	15	5

Table 1: decision tree training data

4. [10 points] Suppose the domain X is the real line, \mathbb{R} , and the labels lie in $Y = \{-1, 1\}$. Let \mathcal{C} be the concept class consisting of simple threshold functions of the form h_θ for some $\theta \in \mathbb{R}$, where $h_\theta(x) = -1$ for all $x \leq \theta$ and $h_\theta(x) = 1$ otherwise. Give a simple and efficient PAC learning algorithm for \mathcal{C} that uses only $m = O(\frac{1}{\epsilon} \log \frac{1}{\delta})$ training examples to output a classifier with error at most ϵ with probability at least $1 - \delta$.
5. [6 points] In this problem we will show that the existence of an efficient mistake-bounded learner for a class \mathcal{C} implies an efficient PAC learner for \mathcal{C} .

Concretely, let \mathcal{C} be a function class with domain $X \in \{-1, 1\}^n$ and binary labels $Y \in \{-1, 1\}$. Assume that \mathcal{C} can be learned by algorithm/learner A with some mistake bound t . You may assume you know the value t . You may also assume that at each iteration, A runs in time polynomial in n and that A only updates its state when it gets an example wrong. The **concrete goal of this problem** is to create a PAC-learning algorithm, B , that can PAC-learn concept class \mathcal{C} with respect to an arbitrary distribution D over $\{-1, 1\}^n$ using algorithm A as a sub-routine.

In order to prove that learner B can PAC-learn concept class \mathcal{C} , we must show that there exists a finite number of examples, m , that we can draw from D such that B produces a hypothesis whose true error is more than ϵ with probability at most δ . First, fix some distribution D on X , and we will assume that the examples are labeled by an unknown $c \in \mathcal{C}$. Additionally, for a hypothesis (i.e. function) $h : X \rightarrow Y$, let $\text{err}(h) = \mathbb{P}_{x \sim D}[h(x) \neq c(x)]$. Formally, we will need to bound m such that the following condition holds:

$$\forall \delta, \epsilon \in [0, 1], \exists m \in \mathbb{N} \mid \mathbb{P}_{x \sim D}[\text{err}(B(\{x\}^m)) > \epsilon] \leq \delta \quad x \sim D \quad (1)$$

where $B(\{x\}^m)$ denotes a hypotheses produced from B with m random draws of x from an arbitrary distribution D .

To find this m , we will first decompose it into blocks of examples of size k and make use of results based on a single block to find the bound necessary for m that satisfies condition 1.

Note: Using the identity $\mathbb{P}[\text{err}(h) > \epsilon] + \mathbb{P}[\text{err}(h) \leq \epsilon] = 1$, we can see that $\mathbb{P}[\text{err}(h) > \epsilon] \leq \delta \Leftrightarrow \mathbb{P}[\text{err}(h) \leq \epsilon] \geq 1 - \delta$, which makes the connection to the definition of PAC-learning discussed in lecture explicit.

- Fix a single arbitrary hypothesis $h' : X \rightarrow Y$ produced by A and determine a lower bound on the number of examples, k , such that $\mathbb{P}[\text{err}(h') > \epsilon] \leq \delta'$. (The contrapositive view would be: with probability at least $1 - \delta'$, it must be the case that $\text{err}(h') \leq \epsilon$. Make sure this makes sense.)
- From part 5a we know that as long as a block is at least of size k , then if that block is classified correctly by a *fixed arbitrary hypothesis* h' we can effectively upper bound the probability of the ‘bad event’ (i.e. A outputs h' s.t. $\text{err}(h') > \epsilon$) by δ' . However, our bound must apply to every h that our algorithm B could output for an arbitrary distribution D over examples. With this in mind, how large should m be so that we can bound all hypotheses that could be output? (You may assume that algorithm B will know the mistake bound throughout the question.)
- Put everything together and fully describe (with proof) a PAC learner that is able to output a hypothesis with a true error at most ϵ with probability at least $1 - \delta$, given a mistake bounded learner A . To do this you should first describe your pseudocode for algorithm B which will use A as a sub-routine (no need for minute details or code, broad

strokes will suffice). Then, prove there exists a finite number of m examples for B to PAC-learn \mathcal{C} for all values of δ and ϵ by lower bounding m by a function of ϵ , δ , and t (i.e. finding a finite lower bound for m such that the PAC-learning requirements in [1](#) are satisfied).