

## Homework 2 - Theory

*Lecture: Prof. Adam Klivans*

*Keywords: Perceptron, SGD, Boosting*

**Instructions:** Please either typeset your answers (L<sup>A</sup>T<sub>E</sub>X recommended) or write them very clearly and legibly and scan them, and upload the PDF on edX. Legibility and clarity are critical for fair grading.

1. **[10 points]** Consider running the Perceptron algorithm on a training set  $S$  arranged in a certain order. Now suppose we run it with the same initial weights and on the *same* training set but in a different order,  $S'$ . Does Perceptron make the same number of mistakes? Does it end up with the same final weights? If so, prove it. If not, give a counterexample, i.e. an  $S$  and  $S'$  where order matters.
2. **[10 points]** We have mainly focused on squared loss, but there are other interesting losses in machine learning. Consider the following loss function which we denote by  $\phi(z) = \max(0, -z)$ . Let  $S$  be a training set  $(x^1, y^1), \dots, (x^m, y^m)$  where each  $x^i \in \mathbb{R}^n$  and  $y^i \in \{-1, 1\}$ . Consider running stochastic gradient descent (SGD) to find a weight vector  $w$  that minimizes  $\frac{1}{m} \sum_{i=1}^m \phi(y^i \cdot w^T x^i)$ . Explain the explicit relationship between this algorithm and the Perceptron algorithm. Recall that for SGD, the update rule when the  $i^{\text{th}}$  example is picked at random is

$$w_{\text{new}} = w_{\text{old}} - \eta \nabla \phi(y^i w^T x^i).$$

*Note:* You do not need to be overly concerned about the discontinuity at  $\phi(0)$ , so you can ignore this when calculating the gradient for this problem.

3. **[6 points]** Here we will give an illustrative example of a weak learner for a simple concept class. Let the domain be the real line,  $\mathbb{R}$ , and let  $\mathcal{C}$  refer to the concept class of “3-piece classifiers”, which are functions of the following form: for  $\theta_1 < \theta_2$  and  $b \in \{-1, 1\}$ ,  $h_{\theta_1, \theta_2, b}(x)$  is  $b$  if  $x \in [\theta_1, \theta_2]$  and  $-b$  otherwise. In other words, they take a certain Boolean value inside a certain interval and the opposite value everywhere else. For example,  $h_{10, 20, 1}(x)$  would be  $+1$  on  $[10, 20]$ , and  $-1$  everywhere else. Let  $\mathcal{H}$  refer to the simpler class of “decision stumps”, i.e. functions  $h_{\theta, b}$  such that  $h(x)$  is  $b$  for all  $x \leq \theta$  and  $-b$  otherwise.
  - (a) Show formally that for any distribution on  $\mathbb{R}$  (assume finite support, for simplicity; i.e., assume the distribution is bounded within  $[-B, B]$  for some large  $B$ ) and any unknown labeling function  $c \in \mathcal{C}$  that is a 3-piece classifier, there exists a decision stump  $h \in \mathcal{H}$  that has error at most  $1/3$ , i.e.  $\mathbb{P}[h(x) \neq c(x)] \leq 1/3$ .
  - (b) Describe a simple, efficient procedure for finding a decision stump that minimizes error with respect to a finite training set of size  $m$ . Such a procedure is called an empirical risk minimizer (ERM).
  - (c) Give a short intuitive explanation for why we should expect that we can easily pick  $m$  sufficiently large that the training error is a good approximation of the true error, i.e. why we can ensure generalization. (Your answer should relate to what we have gained in

going from requiring a learner for  $\mathcal{C}$  to requiring a learner for  $\mathcal{H}$ .) This lets us conclude that we can weakly learn  $\mathcal{C}$  using  $\mathcal{H}$ .

4. **[10 points]** Consider an iteration of the AdaBoost algorithm (using notation from the video lecture on Boosting) where we have obtained classifier  $h_t$ . Show that with respect to the distribution  $D_{t+1}$  generated for the next iteration,  $h_t$  has accuracy exactly  $1/2$ .