

# hw6-programming-solved

November 9, 2024

## 1 Kernel Regression

Given a training dataset  $\{x_i, y_i\}_{i=1}^n$ , kernel regression approximates the unknown nonlinear relation between  $x$  and  $y$  with a function of form

$$y \approx f(x; w) = \sum_{i=1}^n w_i k(x, x_i),$$

where  $k(x, x')$  is a positive definite kernel specified by the users, and  $w_i$  is a set of weights. We will use the simple Gaussian radius basis function (RBF) kernel,

$$k(x, x') = \exp\left(-\frac{\|x - x'\|^2}{2h^2}\right),$$

where  $h$  is a bandwidth parameter.

### 1.0.1 Step 1. Simulate a 1-dimensional dataset


```
[28]: import numpy as np
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline

np.random.seed(100)

### Step 1: Simulate a simple 1D data ###
xTrain = np.expand_dims(np.linspace(-5, 5, 100), 1) # 100*1
yTrain = np.sin(xTrain) + 0.5*np.random.uniform(-1, 1, size=xTrain.shape) ##
↪100 *1

print('xTrain shape', xTrain.shape, 'yTrain shape', yTrain.shape)
plt.plot(xTrain, yTrain, '*')
plt.show()
```

xTrain shape (100, 1) yTrain shape (100, 1)



Now we have a dataset with 100 training data points. Let us calculate the kernel function.

### 1.0.2 Step 2. Kernel function

Your task is to complete the following `rbf_kernel` function that takes two sets of points  $X$  (of size  $n$ ) and  $X'$  (of size  $m$ ) and the bandwidth  $h$  and outputs their pairwise kernel matrix  $K = [k(x_i, x_j)]_{ij}$ , which is of size  $n \times m$ . We will represent input data as matrices, with  $X = [x_i]_{i=1}^n \in \mathbb{R}^{n \times 1}$  denoting the input features and  $Y = [y_i]_{i=1}^n \in \mathbb{R}^{n \times 1}$  the input labels.)

```
[29]: """
        calculating kernel matrix between X and Xp
        """
def rbf_kernel(X, Xp, h):
    # X: n*1 matrix
    # Xp: m*1 matrix
    # h: scalar value

    dists = np.square(X - Xp.T)           # Obtain pairwise differences between X
    ↪and Xp,                               # then square to get pairwise squared
    ↪distances
    K = np.exp(-dists/(2 * h**2))         # Apply the RBF kernel formula
```

```

return K # n*m matrix

### evaluation: if your implementation is correct, you should expect the output
↳ is a 2X3 matrix
# [[0.60653066 1.          0.60653066]
#  [0.13533528 0.60653066 1.          ]]
k_test = rbf_kernel(np.array([[2],[1]]), np.array([[3],[2],[1]]), 1)
print(k_test)

```

```

[[0.60653066 1.          0.60653066]
 [0.13533528 0.60653066 1.          ]]

```

### 1.0.3 Step 3. The median trick for bandwidth

The choice of the bandwidth  $h$  A common way to set the bandwidth  $h$  in practice is the so called median trick, which sets  $h$  to be the median of the pairwise distance on the training data, that is

$$h_{med} = \text{median}(\{\|x_i - x_j\| : i \neq j, \quad i, j = 1, \dots, n\}).$$

- Task: Complete the median distance function.

```

[30]: from scipy.spatial import distance

def median_distance(X):
    # X: n*1 matrix

    #TODO: Calculate the median of the pairwise distance of $X$ below
    # (hint: use '[dist[i, j] for i in range(len(X)) for j in range(len(X)) if i
    ↳ != j]' to remove the diagonal terms; use np.median)

    # Compute pairwise distances using scipy's distance function, resulting in
    ↳ an n x n matrix
    pairwise_distances = distance.cdist(X, X, metric='euclidean')

    # Extract all non-self-distance elements (i.e., i != j)
    off_diagonal_distances = [
        pairwise_distances[i, j]
        for i in range(len(X))
        for j in range(len(X))
        if i != j
    ]

    # Compute the median of these distances
    h = np.median(off_diagonal_distances)

```

```

    return h

### Test your functions
#evaluation: if your implementation is correct, your answer should be [2.0]
h_test = median_distance(np.array([[1],[2],[4]]))
print(h_test)

```

2.0

#### 1.0.4 Step 4. Kernel regression

The weights  $w_i$  are estimated by minimizing a regularized mean square error:

$$\min_w \left( \sum_{i=1}^n (y_i - f(x_i; w))^2 \right) + \beta w^\top K w,$$

where  $w$  is the column vector formed by  $w = [w_i]_{i=1}^n$  and  $K$  is the kernel matrix.

- Please derive the optimal solution of  $w$  using matrix inversion (no need to show the work)
- Complete the following function to implement the calculation of  $w$

```

[31]: def kernel_regression_fitting(xTrain, yTrain, h, beta=1):
    # X: input data, numpy array, n*1
    # Y: input labels, numpy array, n*1

    # Step 1: Compute the kernel matrix K (n x n)
    K = rbf_kernel(xTrain, xTrain, h)

    # Step 2: Add the regularization term to the kernel matrix
    n = K.shape[0]
    K_regularized = K + beta * np.eye(n)

    # Step 3: Compute the weights w using matrix inversion
    W = np.linalg.inv(K_regularized).dot(yTrain)

    return W

### evaluating your code, the shape should be (100, 1) (check the values
↳yourself)
h = median_distance(xTrain)
W_test = kernel_regression_fitting(xTrain, yTrain, h)
print(W_test.shape)

```

(100, 1)

#### 1.0.5 Step 5. Evaluation and Cross Validation

We now need to evaluate the algorithm on the testing data and select the hyperparameters (bandwidth and regularization coefficient) using cross validation

```
[32]: # Please run and read the following base code

def kernel_regression_fit_and_predict(xTrain, yTrain, xTest, h, beta):

    #fitting on the training data
    W = kernel_regression_fitting(xTrain, yTrain, h, beta)

    # computing the kernel matrix between xTrain and xTest
    K_xTrain_xTest = rbf_kernel(xTrain, xTest, h)

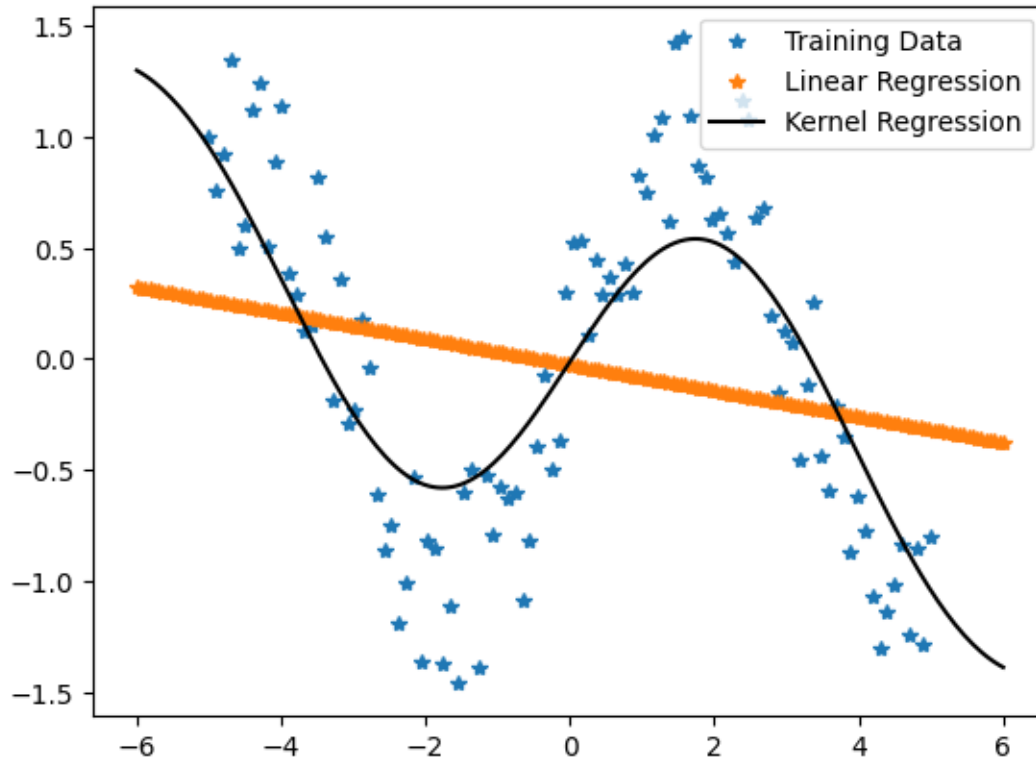
    # predict the label of xTest
    yPred = np.dot( K_xTrain_xTest.T, W)
    return yPred

# generate random testing data
xTest = np.expand_dims(np.linspace(-6, 6, 200), 1) ## 200*1

beta = 1.
# calculating bandwidth
h_med = median_distance(xTrain)
yHatK = kernel_regression_fit_and_predict(xTrain, yTrain, xTest, h_med, beta)

# we also add linear regression for comparision
from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(xTrain, yTrain)
yHat = lr.predict(xTest) # prediction

# visulization
plt.plot(xTrain, yTrain, '*', label='Training Data')
plt.plot(xTest, yHat, '*', label='Linear Regression')
plt.plot(xTest, yHatK, '-k', label='Kernel Regression')
plt.legend()
plt.show()
```



### 1.0.6 Step 5.1. Impact of bandwidth

Run the kernel regression with regularization coefficient  $\beta = 1$  and bandwidth  $h \in \{0.1h_{med}, h_{med}, 10h_{med}\}$ .

- Task: Show the curve learned by different  $h$ . Comment on how  $h$  influences the smoothness of  $h$ .

```
[33]: ### fitting on the training data ###
beta = 1

plt.figure(figsize=(12, 4))

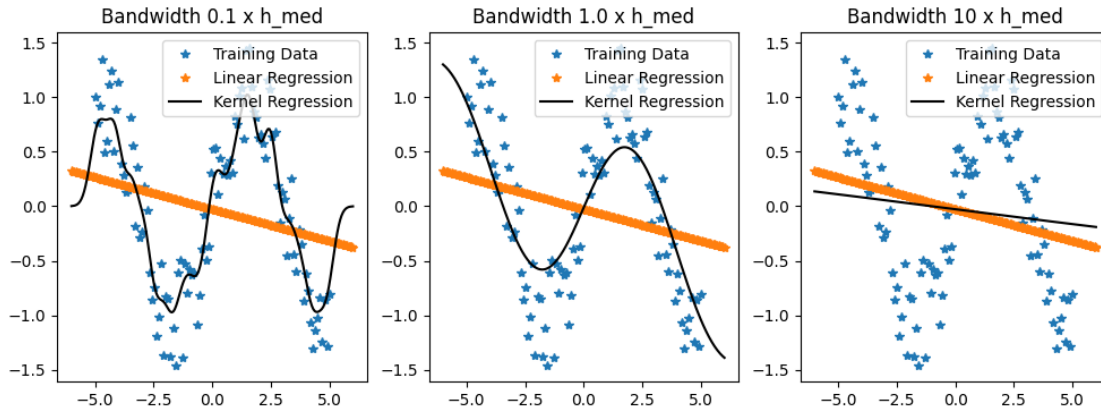
for i, coeff in enumerate([0.1, 1., 10]):
    plt.subplot(1, 3, i+1)

    ### TODO: run kernel regression with bandwidth h = coeff * h_med.
    h = coeff * h_med
    yHat_k_i = kernel_regression_fit_and_predict(xTrain, yTrain, xTest, h, beta)

    # visualization
    plt.plot(xTrain, yTrain, '*', label='Training Data')
    plt.plot(xTest, yHat_k_i, '*', label='Linear Regression')
```

```
plt.plot(xTest, yHatk_i, '-k', label='Kernel Regression')
plt.title('Bandwidth {} x h_med'.format(coff))
plt.legend()
```

```
plt.show()
```



### 1.0.7 How the Bandwidth, $h$ , Influences the Smoothness of the Curve

It is apparent that the higher the value of  $h$ , the smoother the curve will be. When  $h = 0.1 \times h_{med}$ , we can see that the curve is quite rough, but has the advantage of fitting the training data very closely. However, it should be noted that this may lead to overfitting.

### 1.0.8 Step 5.2. Cross Validation (CV)

Use 5-fold cross validation to find the optimal combination of  $h$  and  $\beta$  within  $h \in \{0.1h_{med}, h_{med}, 10h_{med}\}$  and  $\beta \in \{0.1, 1\}$ . - Task: complete the code of cross validation and find the best  $h$  and  $\beta$ . Plot the curve fit with the optimal hyperparameters.

```
[34]: best_beta, best_coff = 1., 1.
best_mse = 1e8

for beta in [0.1, 1]:
    for coff in [0.1, 1., 10.]:
        # 5-fold cross validation
        max_fold = 5
        mse = []

        for i in range(max_fold):

            #TODO: calculate the index of the training/testing partition
            ↪ within 5 fold CV.
            # (hint: set trnIdx to be these index with idx%max_fold!=i, and
            ↪ testIdx with idx%max_fold==i)
```

```

trnIdx = [j for j in range(len(xTrain)) if j % max_fold != i]
testIdx = [j for j in range(len(xTrain)) if j % max_fold == i]

i_xTrain, i_yTrain = xTrain[trnIdx], yTrain[trnIdx]
i_xValid, i_yValid = xTrain[testIdx], yTrain[testIdx]

##TODO: run kernel regression on (i_xTrain, i_yTrain) and calculate
the mean square error on (i_xValid, i_yValid)
h = coff * h_med
i_yPred = kernel_regression_fit_and_predict(i_xTrain, i_yTrain,
i_xValid, h, beta)
mse.append((i_yValid - i_yPred)**2)

mse = np.mean(mse)

# keep track of the combination with the best MSE
if mse < best_mse:
    best_beta, best_coff = beta, coff
    best_mse = mse

print('Beta beta', best_beta, 'Best bandwith', '{}*h_med'.format(best_coff),
    'mse', best_mse)

# bandwidth
h = best_coff * median_distance(xTrain)
yHat_k_i = kernel_regression_fit_and_predict(xTrain, yTrain, xTest, h, best_beta)

# visulization
plt.plot(xTrain, yTrain, '*', label='Training Data')
plt.plot(xTest, yHat, '*', label='Linear Regression')
plt.plot(xTest, yHat_k_i, '-k', label='Kernel Regression (Optimal)')
plt.title('beta {}, bandwidth {}*h_med'.format(best_beta, best_coff))
plt.legend(loc='upper right', fontsize='small')
plt.show()

```

Beta beta 1 Best bandwith 0.1\*h\_med mse 0.1116622935589619



